

# AI Engineer Intern (Agent Systems)

**Location:** Remote-friendly

**Role Type:** Intern (3 months), Individual Contributor

---

## About the Company

PluralHire is building a full-stack talent and career intelligence platform operating across both demand and supply sides of the market.

On the demand side, we help modern organisations build AI-ready, techno-functional GTM teams - specializing in executive hiring, GTM pods, and AI-fluent talent across India, MEA, Europe, and the US.

On the supply side, we are building an end-to-end career transition platform - helping individuals move from role discovery and selection to landing and successfully settling into new roles through diagnostics, structured insights, and personalized guidance.

At the core, PluralHire combines AI-driven diagnostics, human judgment, and market intelligence to solve one of the hardest problems - matching the right talent to the right opportunity, with clarity and precision.

---

## About the Role

This role is focused on the supply-side diagnostic engine, specifically building the AI intelligence layer that powers career assessments.

You will work on multi-agent systems, prompt engineering, and orchestration pipelines - ensuring outputs are accurate, structured, and production-ready.

This is a role for someone who wants to move beyond experimentation and help build reliable AI systems inside a real product.

# What You'll Do

## AI Systems (Core Focus - 60%)

- Work with pipelines of specialized AI agents responsible for resume analysis, skill scoring, motivation assessment, and report generation
- Optimize prompts for accuracy, consistency, and speed
- Improve reliability, reduce hallucinations, and strengthen output quality

## Agent Orchestration (25%)

- Design execution flows across agents and manage dependencies
- Enable retries, validation steps, and structured output enforcement
- Explore where parallel execution can improve speed and throughput

## Evaluation & Quality (15%)

- Build evaluation datasets and benchmarking approaches
  - Ensure every key insight is traceable to real evidence in candidate data
  - Continuously improve output quality and reliability over time
- 

## What do we look for ?

- Strong in Python and/or TypeScript
- Familiar with LLM APIs such as OpenAI, Anthropic, or similar tools
- Understanding of prompt engineering techniques and structured outputs
- Comfortable working with JSON schemas, validations, and pipeline patterns
- Bonus: exposure to LangChain, RAG, evaluation tooling, or agent frameworks

## Who You Are

- A builder who thinks in systems and intelligence layers
  - Curious about AI behaviour, output quality, and production reliability
  - Detail-oriented with a bias for experimentation
  - Interested in real-world AI applications, not just demos
- 

## Why Join Us

We are invested in your long-term trajectory.

This role gives you direct exposure to how production-grade AI systems are designed, evaluated, and improved. You will work on real multi-agent architecture and learn how AI moves from prompt to product.

---

## Application Details

Share your CV + portfolio at:

[pulkit@pluralhire.com](mailto:pulkit@pluralhire.com)

[hullas@pluralhire.com](mailto:hullas@pluralhire.com)

[arora.hullas@gmail.com](mailto:arora.hullas@gmail.com)